

NLP: Yelp Review to Rating

Authors: Tanvee Desai and Tanner Arrizabalaga

Hello! In this project, we will be looking over Yelp reviews (data available here: <https://www.yelp.com/dataset> (<https://www.yelp.com/dataset>)) and utilizing ML/DL to accurately predict what the reviews star rating is based solely on text.

This project is split into the following parts

- Libraries
- EDA
- Data Cleaning
 - Stop word removal, HTML parsing, punctuation removal, etc.
 - Creation of a cleaned *and* stemmed dataset
- Model Implementation
 - Simple BOW Model Neural Network
 - LSTM
 - One vs. All LSTM Approach
- Exploring Challenges
 - Challenge 5
 - Challenge 6

Importing necessary libraries

```
In [ ]: # General Libraries
import json
import sys
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import itertools

# NLP
import nltk
import re
from nltk.corpus import stopwords
from bs4 import BeautifulSoup
from nltk.stem import PorterStemmer

# ML/DL
import tensorflow as tf
import pickle

from sklearn.preprocessing import LabelBinarizer, LabelEncoder
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split

from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Activation, Dropout, Embedding, Conv1D, MaxPooling1D, LSTM, BatchNormalization, SpatialDropout1D, Bidirectional
from keras.preprocessing.sequence import pad_sequences
from keras.preprocessing import text, sequence
from keras import utils
from keras import regularizers
from keras.models import load_model
```

```
In [ ]: yelp = pd.read_json("./yelp_review_training_dataset.jsonl", lines = True)
yelp.head()
```

How large is the data?

```
In [ ]: yelp.shape
```

EDA - Stars

Not too much to go off of, but let's get a general understanding of our data. How many nulls do we have?

```
In [ ]: yelp.isna().sum()
```

```
In [ ]: sns.countplot(yelp['stars'])
```

One thing we can potentially look at is whether or not the reviews are balanced. Let's say ≥ 4 is positive, and < 4 is negative. If we do see a significant difference in positive and negative reviews, we can balance it before training.

```
In [ ]: def pos_or_neg(x):  
        if x >= 4:  
            return "Positive"  
        else:  
            return "Negative"  
  
yelp['category'] = yelp['stars'].apply(pos_or_neg)  
  
sns.countplot(yelp['category'])  
num_pos = np.count_nonzero(yelp['category'] == 'Positive')  
num_neg = np.count_nonzero(yelp['category'] == 'Negative')  
print("Positive to negative review ratio: ", num_pos / num_neg)
```

There are roughly 1 and 2/3 times as many positive reviews as negative reviews. We will first try no class balancing when building the model, but may turn to class balancing later on.

Data Cleaning - Text

```

In [ ]: REPLACE_BY_SPACE_RE = re.compile('[/(){}\\[\\]\\|@,;]')
BAD_SYMBOLS_RE = re.compile('[^0-9a-z #+_]')
STOPWORDS = set(stopwords.words('english'))
print(STOPWORDS)

def adjust_stopwords(stopwords):
    words_to_keep = set(['nor', 'not', 'very', 'no', 'few', 'too', 'doesn', 'd
    idn', 'wasn', 'ain',
                        "doesn't", "isn't", "hasn't", 'shouldn', "weren't", "d
    on't", "didn't",
                        "shouldn't", "wouldn't", "won't", "above", "below", "h
    aven't", "shan't", "weren"])
    return stopwords - words_to_keep

def clean_text(text):
    """
    text: a string

    return: modified initial string
    """
    new_text = BeautifulSoup(text, "lxml").text # HTML decoding
    new_text = new_text.lower() # lowercase text
    new_text = REPLACE_BY_SPACE_RE.sub(' ', new_text) # replace REPLACE_BY_SPACE_RE symbols by space in text
    new_text = BAD_SYMBOLS_RE.sub(' ', new_text) # delete symbols which are in BAD_SYMBOLS_RE from text

    ps = PorterStemmer()

    new_text = ' '.join(ps.stem(word) for word in new_text.split()) # keeping all words, no stop word removal
    # new_text = ' '.join(ps.stem(word) for word in new_text.split() if word not in STOPWORDS) # delete stopwords from text and stem
    return new_text

STOPWORDS = adjust_stopwords(STOPWORDS)
print(STOPWORDS)

```

```

In [ ]: %%time
yelp['text'] = yelp['text'].apply(clean_text)
yelp.to_csv('cleaned_yelp_stemmed.csv')

```

```
In [ ]: text_1 = "\"Good morning, cocktails for you?\" \"Wait...what? Oh...it's Vegas!
\n\nDining here, you best not be dieting because this place is literally the d
efinition of excess, but in a good way. I'm a sucker for benedicts so that was
awesome. \"Service was really great too and the staff was so welcoming. It was
our first stop just after landing so really appreciate the service.\n\nBack in
Hawaii this reminds me of Zippys or Anna Millers - that home feeling. Prices a
re a bit high, but for what you get it's totally worth it. Will remember this
place if I ever return to Vegas in the future.\"
text_2 = \"80 bucks, thirty minutes to fix my shattered iPhone screen. Verizon
won't help you so go here\"
text_3 = \"Tr\u00e8s grand caf\u00e9, mais aussi calme et reposant, je m'y suis
arr\u00eat\u00e9 alors que j'\u00e9tais dans le coin.\n\nOn peu y mang\u00e9 l
e midi, prendre une p\u00e2tisserie ou un caf\u00e9/th\u00e9. \"J'ai prit un
th\u00e9 qui \u00e9tait vraiment bon, et je me suis pos\u00e9 devant une des g
randes baies vitr\u00e9es sur un coussin et j'ai relax\u00e9 compl\u00e8tement
pendant 2 heures. \"Mais c'est aussi une coop\u00e9rative d'artiste, avec un
e estrade etc.\n\nIl y a aussi un magasin Bio \u00e0 l'entr\u00e9e o\u00f9 vou
s retrouverez des savons, huile d'olive et plein d'autres produits.\"
text_4 = \"Sadly, as of July 28, 2016, Silverstein bakery is permanently close
d. I went there today in person and found the bad news posted on their door. :
(\"
text_5 = \"I went here they were about to close but the cashier was especially
helpful ..but I guess they were tired of work...\"

clean_text(text_1)
```

Model Implementation

Evaluation

1. Average Star Error (Average Absolute offset between predicted and true number of stars)
2. Accuracy (Exact Match -- Number of exactly predicted star ratings / total samples)

```
In [49]: def MAE(y_true, y_pred):
    diffs = np.abs(y_true - y_pred)
    loss = np.mean(diffs)
    return loss

def Accuracy(y_true, y_pred):
    correct = y_true == y_pred
    cor_count = np.count_nonzero(correct)
    return cor_count / len(y_true)
```

Train/Test Split (Unbalanced and balanced)

```
In [50]: yelp = pd.read_csv('cleaned_yelp_stemmed.csv')
yelp.head()
```

Out[50]:

	Unnamed: 0	review_id	text	stars	category
0	0	Q1sbwVQXV2734tPgoKj4Q	total bill for thi horribl servic over 8g thes...	1	Negative
1	1	GJXCdrto3ASJOqKeVWPi6Q	i ador travi at the hard rock s new kelli card...	5	Positive
2	2	2TzJjDVDEuAW6MR5Vuc1ug	i have to say that thi offic realli ha it toge...	5	Positive
3	3	yi0R0Ugj_xUx_Nek0-_Qig	went in for a lunch steak sandwich wa delici a...	5	Positive
4	4	11a8sVPMUFtaC7_ABRkmtw	today wa my second out of three session i had ...	1	Negative

```
In [51]: X = yelp['text'].fillna('').values
y = yelp['stars']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)
```

Baseline Sequential Model

```
In [52]: max_words = 3000
tokenizer = text.Tokenizer(num_words=max_words, char_level=False)

tokenizer.fit_on_texts(X_train)
X_train = tokenizer.texts_to_matrix(X_train)
X_test = tokenizer.texts_to_matrix(X_test)

encoder = LabelEncoder()
encoder.fit(y_train)
y_train = encoder.transform(y_train)
y_test = encoder.transform(y_test)

num_classes = np.max(y_train) + 1
y_train = utils.to_categorical(y_train, num_classes)
y_test = utils.to_categorical(y_test, num_classes)

print('X_train shape:', X_train.shape)
print('X_test shape:', X_test.shape)
print('y_train shape:', y_train.shape)
print('y_test shape:', y_test.shape)

X_train shape: (373506, 3000)
X_test shape: (160075, 3000)
y_train shape: (373506, 5)
y_test shape: (160075, 5)
```

Let's save the tokenizer as well for our test submission file script.

```
In [ ]: # saving
with open('tokenizer.pickle', 'wb') as handle:
    pickle.dump(tokenizer, handle, protocol=pickle.HIGHEST_PROTOCOL)

# Loading
with open('tokenizer.pickle', 'rb') as handle:
    tokenizer = pickle.load(handle)
```

Here, we are computing a single model, but in future we will optimize on several parameters, listed below

- Batch size
- Learning rate
- Gradient clipping
- Drop out
- Batch normalization
- Optimizers
- Regularization

After some tests, the main variations I noticed were from the learning rate, regularization, and the choice of the optimizer. With that being said, this baseline model will use **ADAM with a learning rate of .0001 and regularization (kernel, bias, and activity)**

```
In [53]: batch_size = 512
epochs = 10

lr_schedule = keras.optimizers.schedules.ExponentialDecay(
    initial_learning_rate=.0001,
    decay_steps=10000,
    decay_rate=0.9)

optimizer = keras.optimizers.Adam(learning_rate=lr_schedule, beta_1=0.9, beta_
2=0.95, amsgrad=False)

baseline = Sequential()
baseline.add(Dense(512, input_shape=(max_words,), kernel_regularizer=regulariz
ers.l1_l2(l1=1e-5, l2=1e-4),
            bias_regularizer=regularizers.l2(1e-4),
            activity_regularizer=regularizers.l2(1e-5)))
baseline.add(BatchNormalization())
baseline.add(Activation('relu'))
baseline.add(Dropout(0.3))
baseline.add(Dense(5))
baseline.add(Activation('softmax'))

baseline.compile(loss='categorical_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

history = baseline.fit(X_train, y_train,
                      batch_size=batch_size,
                      epochs=epochs,
                      verbose=1,
                      validation_split=0.2)
```


Train on 298804 samples, validate on 74702 samples

Epoch 1/10

298804/298804 [=====] - 22s 73us/step - loss: 1.3055
- accuracy: 0.6878 - val_loss: 1.1030 - val_accuracy: 0.7437

Epoch 2/10

298804/298804 [=====] - 12s 40us/step - loss: 1.0671
- accuracy: 0.7486 - val_loss: 1.0346 - val_accuracy: 0.7491

Epoch 3/10

298804/298804 [=====] - 12s 39us/step - loss: 0.9772
- accuracy: 0.7633 - val_loss: 0.9830 - val_accuracy: 0.7512

Epoch 4/10

298804/298804 [=====] - 12s 39us/step - loss: 0.9067
- accuracy: 0.7749 - val_loss: 0.9426 - val_accuracy: 0.7507

Epoch 5/10

298804/298804 [=====] - 12s 39us/step - loss: 0.8487
- accuracy: 0.7843 - val_loss: 0.9106 - val_accuracy: 0.7525

Epoch 6/10

298804/298804 [=====] - 11s 38us/step - loss: 0.8001
- accuracy: 0.7921 - val_loss: 0.8874 - val_accuracy: 0.7518

Epoch 7/10

298804/298804 [=====] - 12s 39us/step - loss: 0.7584
- accuracy: 0.8002 - val_loss: 0.8670 - val_accuracy: 0.7508

Epoch 8/10

298804/298804 [=====] - 12s 39us/step - loss: 0.7230
- accuracy: 0.8077 - val_loss: 0.8542 - val_accuracy: 0.7490

Epoch 9/10

298804/298804 [=====] - 12s 39us/step - loss: 0.6899
- accuracy: 0.8165 - val_loss: 0.8438 - val_accuracy: 0.7509


Epoch 10/10

298804/298804 [=====] - 12s 39us/step - loss: 0.6618
- accuracy: 0.8232 - val_loss: 0.8362 - val_accuracy: 0.7511


```
In [54]: score = baseline.evaluate(X_test, y_test,
                                   batch_size=batch_size, verbose=1)
print('Test accuracy:', score[1])
```

160075/160075 [=====] - 13s 80us/step
Test accuracy: 0.7536966800689697

```
In [55]: plt.title('Loss')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='test')
plt.legend()
plt.show()
```



```
In [56]: plt.title('Accuracy')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='test')
plt.legend()
plt.show()
```



```
In [ ]: baseline.save('./models/baseline.h5')
```

Now training with several parameter changes

```
In [ ]: batch_sizes = [128, 256, 512]
epochs = [5]
learning_rates = [.01, .001, .0001]
dropout = [False, True]
batch_norm = [False, True]
regularization = [True]
optimizers = ["SGD", "RMSProp", "ADAM"]

all_lists = [batch_sizes, epochs, learning_rates, dropout, batch_norm, regularization, optimizers]

params_to_test = list(itertools.product(*all_lists))
print(len(params_to_test))
```

```
In [ ]: models = {}
        histories = {}
        scores = {}

        for params in params_to_test:
            print(params)
            batch_size, epochs, learning_rate, dropout, batch_norm, regularization, opt = params

            if opt == "SGD":
                optimizer = keras.optimizers.SGD(learning_rate=learning_rate, momentum=0.0, nesterov=False)
            elif opt == "RMSProp":
                optimizer = keras.optimizers.RMSprop(learning_rate=learning_rate, rho=0.9)
            elif opt == "ADAM":
                optimizer = keras.optimizers.Adam(learning_rate=learning_rate, beta_1=0.9, beta_2=0.99, amsgrad=False)
            else:
                optimizer = keras.optimizers.Adadelta(learning_rate=learning_rate, rho=0.95)

            model = Sequential()
            model.add(Dense(512, input_shape=(max_words,), kernel_regularizer=regularizers.l1_l2(l1=1e-5, l2=1e-4)))

            # Check Batch Normalization
            if batch_norm:
                model.add(BatchNormalization())

            model.add(Activation('relu'))

            # Check Dropout
            if dropout:
                model.add(Dropout(0.2))

            model.add(Dense(5))
            model.add(Activation('softmax'))

            model.compile(loss='categorical_crossentropy',
                          optimizer=optimizer,
                          metrics=['accuracy'])

            history = model.fit(X_train, y_train,
                               batch_size=batch_size,
                               epochs=epochs,
                               verbose=0,
                               validation_split=0.1)

            models[params] = model
            histories[params] = history

            score = model.evaluate(X_test, y_test, batch_size=batch_size, verbose=1)
            print(score)

            scores[params] = score
```

LSTM Model

Specific Data Prep

```
In [33]: X = yelp['text'].fillna('').values
y = pd.get_dummies(yelp['stars']).values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)

max_words = 3000
maxlen = 400

X_train = tokenizer.texts_to_sequences(X_train)
X_test = tokenizer.texts_to_sequences(X_test)

# For the LSTM, we are going to pad our sequences
X_train = pad_sequences(X_train, maxlen=maxlen)
X_test = pad_sequences(X_test, maxlen=maxlen)

(373506,) (373506, 5)
(160075,) (160075, 5)
```

LSTM #1

```
In [34]: batch_size = 512
epochs = 10

lr_schedule = keras.optimizers.schedules.ExponentialDecay(
    initial_learning_rate=.001,
    decay_steps=10000,
    decay_rate=0.9)

optimizer = keras.optimizers.Adam(learning_rate=lr_schedule, beta_1=0.9, beta_2=0.99, amsgrad=False, clipvalue=.3)

lstm = Sequential()
lstm.add(Embedding(max_words, 128, input_length=maxlen))
lstm.add(SpatialDropout1D(0.2))
lstm.add(Conv1D(64, 5, activation='relu', kernel_regularizer=regularizers.l1_l2(l1=1e-5, l2=1e-4),
    bias_regularizer=regularizers.l2(1e-4)))
lstm.add(MaxPooling1D(pool_size=4))
lstm.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
lstm.add(BatchNormalization())
lstm.add(Dense(5, activation='sigmoid'))

lstm.compile(loss='categorical_crossentropy',
    optimizer=optimizer,
    metrics=['accuracy'])

history = lstm.fit(X_train, y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_split=0.2)
```

Train on 298804 samples, validate on 74702 samples

Epoch 1/10

298804/298804 [=====] - 87s 292us/step - loss: 0.7897 - accuracy: 0.7041 - val_loss: 0.6659 - val_accuracy: 0.7420

Epoch 2/10

298804/298804 [=====] - 87s 292us/step - loss: 0.6497 - accuracy: 0.7528 - val_loss: 0.6169 - val_accuracy: 0.7640

Epoch 3/10

298804/298804 [=====] - 85s 285us/step - loss: 0.6091 - accuracy: 0.7650 - val_loss: 0.6040 - val_accuracy: 0.7672

Epoch 4/10

298804/298804 [=====] - 85s 285us/step - loss: 0.5879 - accuracy: 0.7719 - val_loss: 0.5929 - val_accuracy: 0.7698

Epoch 5/10

298804/298804 [=====] - 85s 286us/step - loss: 0.5724 - accuracy: 0.7776 - val_loss: 0.5925 - val_accuracy: 0.7728

Epoch 6/10

298804/298804 [=====] - 85s 284us/step - loss: 0.5603 - accuracy: 0.7828 - val_loss: 0.5810 - val_accuracy: 0.7765

Epoch 7/10

298804/298804 [=====] - 86s 287us/step - loss: 0.5490 - accuracy: 0.7872 - val_loss: 0.5755 - val_accuracy: 0.7785

Epoch 8/10

298804/298804 [=====] - 85s 284us/step - loss: 0.5391 - accuracy: 0.7906 - val_loss: 0.5791 - val_accuracy: 0.7777

Epoch 9/10

298804/298804 [=====] - 85s 284us/step - loss: 0.5311 - accuracy: 0.7942 - val_loss: 0.5778 - val_accuracy: 0.7806

Epoch 10/10

298804/298804 [=====] - 85s 284us/step - loss: 0.5229 - accuracy: 0.7981 - val_loss: 0.5802 - val_accuracy: 0.7784

LSTM #1: Evaluation

```
In [35]: score = lstm.evaluate(X_test, y_test,
                                batch_size=batch_size, verbose=1)
print('Test accuracy:', score[1])
```

160075/160075 [=====] - 11s 69us/step


Test accuracy: 0.7798094749450684

In [36]: `lstm.summary()`


Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 400, 128)	384000
spatial_dropout1d_2 (Spatial	(None, 400, 128)	0
conv1d_2 (Conv1D)	(None, 396, 64)	41024
max_pooling1d_2 (MaxPooling1	(None, 99, 64)	0
lstm_2 (LSTM)	(None, 128)	98816
batch_normalization_4 (Batch	(None, 128)	512
dense_6 (Dense)	(None, 5)	645
Total params: 524,997		
Trainable params: 524,741		
Non-trainable params: 256		

In [37]: `plt.title('Loss')`
`plt.plot(history.history['loss'], label='train')`
`plt.plot(history.history['val_loss'], label='test')`
`plt.legend()`
`plt.show()`




```
In [38]: plt.title('Accuracy')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='test')
plt.legend()
plt.show()
```



Let's save this model as well.

```
In [ ]: lstm.save('./models/lstm.h5')
```

One vs. All Approach

In the one vs. all approach, it goes by the following idea:

- We will have N learners for the multi-class classification problem, where N is the number of classes
- For each learner L , we will train L on our training data X_{Train} and y_{Train} . However, y_{Train} consists of only one label, making it a binary classification problem instead of multinomial
 - For instance, learner L_1 will still use all of X_{Train} , but y_{Train} will now be transformed to be a binary vector v_i where i denotes the star rating we are attempting to predict
- Once we have concluded our training, we will then create an ensemble model (bagging) that does the following
 1. L_1, L_2, \dots, L_5 all assign p_i to each record in X_{Test} , where p_i is the likelihood observation x_n belongs to class i
 2. From there, our prediction is the following: $P_n = \text{argmax}(p_1, p_2, p_3, p_4, p_5)$

After observing the challenge datasets 5 & 6, my partner and I believe this approach is a clever way to tackle the challenges while still having a strong model.

Sources: <https://developers.google.com/machine-learning/crash-course/multi-class-neural-networks/one-vs-all>
 (<https://developers.google.com/machine-learning/crash-course/multi-class-neural-networks/one-vs-all>)

```
In [39]: yelp = pd.read_csv('cleaned_yelp_stemmed.csv')

X = yelp['text'].fillna('').values
y = pd.get_dummies(yelp['stars']).values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)

# Loading
with open('tokenizer.pickle', 'rb') as handle:
    tokenizer = pickle.load(handle)

max_words = 3000
maxlen = 400

X_train = tokenizer.texts_to_sequences(X_train)
X_test = tokenizer.texts_to_sequences(X_test)
X_train = pad_sequences(X_train, maxlen=maxlen)
X_test = pad_sequences(X_test, maxlen=maxlen)

print('X_train shape:', X_train.shape)
print('X_test shape:', X_test.shape)
print('y_train shape:', y_train.shape)
print('y_test shape:', y_test.shape)

X_train shape: (373506, 400)
X_test shape: (160075, 400)
y_train shape: (373506, 5)
y_test shape: (160075, 5)
```

Buidling all models

```

In [40]: stars = np.arange(1, 6)
models = {}
histories = {}
batch_size = 1024

for star in stars:
    if star in [1]:
        epochs = 2
    elif star in [2, 3, 4]:
        epochs = 3
    else:
        epochs = 4

    print(star)
    y_train_sub = y_train[:, star - 1]

    lr_schedule = keras.optimizers.schedules.ExponentialDecay(
        initial_learning_rate=.001,
        decay_steps=10000,
        decay_rate=0.9)

    optimizer = keras.optimizers.Adam(learning_rate=lr_schedule, beta_1=0.9, b
eta_2=0.99, amsgrad=False, clipvalue=.3)

    sub_lstm = Sequential()
    sub_lstm.add(Embedding(max_words, 128, input_length=maxlen))
    sub_lstm.add(SpatialDropout1D(0.2))
    sub_lstm.add(Conv1D(64, 5, activation='relu', kernel_regularizer=regulariz
ers.l1_l2(l1=1e-5, l2=1e-4),
                bias_regularizer=regularizers.l2(1e-4)))
    sub_lstm.add(MaxPooling1D(pool_size=4))
    sub_lstm.add(LSTM(128))
    sub_lstm.add(BatchNormalization())
    sub_lstm.add(Dense(8))
    sub_lstm.add(Dense(1, activation='sigmoid'))

    sub_lstm.compile(loss='binary_crossentropy',
                    optimizer=optimizer,
                    metrics=['accuracy'])

    history = sub_lstm.fit(X_train, y_train_sub,
                          batch_size=batch_size,
                          epochs=epochs,
                          verbose=1,
                          validation_split=0.2)

    models[star] = sub_lstm
    histories[star] = sub_lstm

```

```
1
Train on 298804 samples, validate on 74702 samples
Epoch 1/3
298804/298804 [=====] - 67s 226us/step - loss: 0.250
9 - accuracy: 0.9039 - val_loss: 0.5960 - val_accuracy: 0.7587
Epoch 2/3
298804/298804 [=====] - 66s 222us/step - loss: 0.178
2 - accuracy: 0.9334 - val_loss: 0.2453 - val_accuracy: 0.8876
Epoch 3/3
298804/298804 [=====] - 66s 220us/step - loss: 0.156
4 - accuracy: 0.9423 - val_loss: 0.3240 - val_accuracy: 0.8869
2
Train on 298804 samples, validate on 74702 samples
Epoch 1/3
298804/298804 [=====] - 67s 224us/step - loss: 0.256
8 - accuracy: 0.9182 - val_loss: 0.2734 - val_accuracy: 0.9323
Epoch 2/3
298804/298804 [=====] - 66s 221us/step - loss: 0.182
3 - accuracy: 0.9353 - val_loss: 0.2096 - val_accuracy: 0.9325
Epoch 3/3
298804/298804 [=====] - 66s 220us/step - loss: 0.163
2 - accuracy: 0.9402 - val_loss: 0.2037 - val_accuracy: 0.9342
3
Train on 298804 samples, validate on 74702 samples
Epoch 1/3
298804/298804 [=====] - 67s 224us/step - loss: 0.251
0 - accuracy: 0.9214 - val_loss: 0.2580 - val_accuracy: 0.9363
Epoch 2/3
298804/298804 [=====] - 66s 221us/step - loss: 0.176
4 - accuracy: 0.9391 - val_loss: 0.2564 - val_accuracy: 0.9363
Epoch 3/3
298804/298804 [=====] - 66s 222us/step - loss: 0.157
9 - accuracy: 0.9444 - val_loss: 0.1955 - val_accuracy: 0.9390
4
Train on 298804 samples, validate on 74702 samples
Epoch 1/3
298804/298804 [=====] - 67s 224us/step - loss: 0.360
6 - accuracy: 0.8547 - val_loss: 0.4197 - val_accuracy: 0.8639
Epoch 2/3
298804/298804 [=====] - 66s 223us/step - loss: 0.304
1 - accuracy: 0.8743 - val_loss: 0.3464 - val_accuracy: 0.8640
Epoch 3/3
298804/298804 [=====] - 66s 221us/step - loss: 0.283
0 - accuracy: 0.8831 - val_loss: 0.3158 - val_accuracy: 0.8696
5
Train on 298804 samples, validate on 74702 samples
Epoch 1/3
298804/298804 [=====] - 67s 224us/step - loss: 0.347
0 - accuracy: 0.8598 - val_loss: 0.5140 - val_accuracy: 0.7406
Epoch 2/3
298804/298804 [=====] - 67s 223us/step - loss: 0.287
5 - accuracy: 0.8851 - val_loss: 0.3091 - val_accuracy: 0.8734
Epoch 3/3
298804/298804 [=====] - 66s 221us/step - loss: 0.260
3 - accuracy: 0.8973 - val_loss: 0.2885 - val_accuracy: 0.8816
```

Building an ensemble model (maximization between learners) for all trained models

Testing

```
In [41]: %%time
# Evaluating the models above (TEST)
y_test_und = pd.DataFrame(y_test)
y_test_true = pd.DataFrame(y_test_und.columns[np.where(y_test_und!=0)[1]]) + 1

# Unload models
lstm_1, lstm_2, lstm_3, lstm_4, lstm_5 = models[1], models[2], models[3], models[4], models[5]

## Predicting the probability for each observation each model
print("Predicting 1 star")
one_star_ps = lstm_1.predict(X_test)
print("Predicting 2 star")
two_star_ps = lstm_2.predict(X_test)
print("Predicting 3 star")
three_star_ps = lstm_3.predict(X_test)
print("Predicting 4 star")
four_star_ps = lstm_4.predict(X_test)
print("Predicting 5 star")
five_star_ps = lstm_5.predict(X_test)

data = [one_star_ps.flatten(), two_star_ps.flatten(), three_star_ps.flatten(),
four_star_ps.flatten(), five_star_ps.flatten()]
cols = [1, 2, 3, 4, 5]
ps = pd.DataFrame(data=data, index=cols).T

ps["pred"] = ps.idxmax(axis=1)
ps.head()

print(MAE(ps["pred"], y_test_true[0]))
print(Accuracy(ps["pred"], y_test_true[0]))
```

```
Predicting 1 star
Predicting 2 star
Predicting 3 star
Predicting 4 star
Predicting 5 star
0.38773699828205527
0.7536592222395752
Wall time: 5min 50s
```

Saving the models

```
In [ ]: # lstm_1.save("./models/one_star.h5")
# lstm_2.save("./models/two_star.h5")
# lstm_3.save("./models/three_star.h5")
# lstm_4.save("./models/four_star.h5")
# lstm_5.save("./models/five_star.h5")
```

Ensemble on Test Set

```
In [42]: yelp = pd.read_csv('cleaned_yelp_stemmed.csv')

X = yelp['text'].fillna('').values
y = pd.get_dummies(yelp['stars'])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)

max_words = 3000
maxlen = 400

# with open('tokenizer.pickle', 'rb') as handle:
#     tokenizer = pickle.load(handle)

print(y_test)

necc_cols = [1, 2, 3, 4, 5]
for col in necc_cols:
    if col not in y_test.columns:
        y_test[col] = 0

y_test = y_test[necc_cols]
y_test = y_test.values

X_baseline = tokenizer.texts_to_matrix(X_test)
X_lstm = tokenizer.texts_to_sequences(X_test)
X_lstm = pad_sequences(X_lstm, maxlen=maxlen)

(373506,) (373506, 5)
(160075,) (160075, 5)
      1  2  3  4  5
255947  0  0  0  0  1
261035  0  0  0  0  1
355633  0  0  0  0  1
205506  0  0  0  0  1
97222   0  0  0  1  0
...     .. .. .. .. ..
491832  0  0  0  0  1
311959  0  0  0  0  1
140524  1  0  0  0  0
125037  0  0  1  0  0
200135  0  0  0  1  0

[160075 rows x 5 columns]
```

```
In [ ]: ## Trying our pretrained models
## Optimizer
# lr_schedule = keras.optimizers.schedules.ExponentialDecay(initial_learning_r
ate=.001, decay_steps=10000, decay_rate=0.9)
# optimizer = keras.optimizers.Adam(learning_rate=lr_schedule, beta_1=0.9, bet
a_2=0.99, amsgrad=False, clipvalue=.3)

## Baseline
# baseline = load_model('./models/baseline.h5')

# baseline.compile(loss='categorical_crossentropy',
#                   optimizer=optimizer,
#                   metrics=['accuracy'])

## LSTM
# lstm = load_model('./models/lstm.h5')

# lstm.compile(loss='categorical_crossentropy',
#              optimizer=optimizer,
#              metrics=['accuracy'])

## One vs. all
# lstm_1 = load_model('./models/one_star.h5')

# lstm_1.compile(loss='binary_crossentropy',
#                optimizer=optimizer,
#                metrics=['accuracy'])

# lstm_2 = load_model('./models/two_star.h5')

# lstm_2.compile(loss='binary_crossentropy',
#                optimizer=optimizer,
#                metrics=['accuracy'])

# lstm_3 = load_model('./models/three_star.h5')

# lstm_3.compile(loss='binary_crossentropy',
#                optimizer=optimizer,
#                metrics=['accuracy'])

# lstm_4 = load_model('./models/four_star.h5')

# lstm_4.compile(loss='binary_crossentropy',
#                optimizer=optimizer,
#                metrics=['accuracy'])

# lstm_5 = load_model('./models/five_star.h5')

# lstm_5.compile(loss='binary_crossentropy',
#                optimizer=optimizer,
#                metrics=['accuracy'])
```

```

In [43]: cols = [1, 2, 3, 4, 5]
# Baseline
print("Baseline")
baseline_preds = pd.DataFrame(baseline.predict(X_baseline), columns=cols)
baseline_preds['baseline_pred'] = baseline_preds.idxmax(axis=1)

# LSTM
print("LSTM")
lstm_preds = pd.DataFrame(lstm.predict(X_lstm), columns=cols)
lstm_preds['lstm_pred'] = lstm_preds.idxmax(axis=1)

# One vs. all
print("OVA")
one_star_ps = lstm_1.predict(X_lstm)
two_star_ps = lstm_2.predict(X_lstm)
three_star_ps = lstm_3.predict(X_lstm)
four_star_ps = lstm_4.predict(X_lstm)
five_star_ps = lstm_5.predict(X_lstm)

data = [one_star_ps.flatten(), two_star_ps.flatten(), three_star_ps.flatten(),
four_star_ps.flatten(), five_star_ps.flatten()]
ova_preds = pd.DataFrame(data=data, index=cols).T

ova_preds["ova_pred"] = ova_preds.idxmax(axis=1)

all_preds = pd.DataFrame([baseline_preds['baseline_pred'], lstm_preds['lstm_pred'],
ova_preds['ova_pred']]).T
all_preds["final_pred"] = all_preds.mode(axis=1)[0]

```

Baseline
LSTM
OVA

```

In [44]: print([MAE(all_preds["final_pred"], pd.DataFrame(data=y_test, columns=cols).idxmax(axis=1)), Accuracy(all_preds["final_pred"], pd.DataFrame(data=y_test, columns=cols).idxmax(axis=1))])

[0.33215055442761204, 0.7702264563485866]

```

Challenges

Challenge 5


```
In [45]: c5 = pd.read_json("./yelp_challenge_5_with_answers.jsonl", lines = True)
print(c5.shape)
c5.head()
```

(500, 3)


Out[45]:

	review_id	text	stars
0	50	I went to this campus for 1 semester. I was in...	2
1	51	I have rated it a two star based on its compar...	2
2	52	Just like most of the reviews, we ordered and ...	2
3	53	I only go here if it is an emergency. I HATE i...	2
4	54	Rude staff. I got 60 feeder fish and about 15 ...	2

Quick EDA

```
In [46]: sns.countplot(c5['stars'])
```

Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x238037a3508>



Pre-processing

```
In [47]: c5['text'] = c5['text'].apply(clean_text)
c5.head()
```

Out[47]:

	review_id	text	stars
0	50	i went to thi campu for 1 semest i wa in busi ...	2
1	51	i have rate it a two star base on it compariso...	2
2	52	just like most of the review we order and paid...	2
3	53	i onli go here if it is an emerg i hate it tha...	2
4	54	rude staff i got 60 feeder fish and about 15 w...	2

Load previous tokenizer

```
In [58]: X = c5['text'].fillna('').values
y = pd.get_dummies(c5['stars'])

# with open('tokenizer.pickle', 'rb') as handle:
#     tokenizer = pickle.load(handle)

max_words

necc_cols = [1, 2, 3, 4, 5]
for col in necc_cols:
    if col not in y.columns:
        y[col] = 0

y = y[necc_cols]
y = y.values

X_baseline = tokenizer.texts_to_matrix(X)
X_lstm = tokenizer.texts_to_sequences(X)
X_lstm = pad_sequences(X_lstm, maxlen=400)
```

Load and compile models

```
In [ ]: # Baseline
baseline = load_model('./models/baseline.h5')

baseline.compile(loss='categorical_crossentropy',
                 optimizer=optimizer,
                 metrics=['accuracy'])

# LSTM
lstm = load_model('./models/lstm.h5')

lstm.compile(loss='categorical_crossentropy',
             optimizer=optimizer,
             metrics=['accuracy'])

# One vs. all
lstm_1 = load_model('./models/one_star.h5')

lstm_1.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_2 = load_model('./models/two_star.h5')

lstm_2.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_3 = load_model('./models/three_star.h5')

lstm_3.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_4 = load_model('./models/four_star.h5')

lstm_4.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_5 = load_model('./models/five_star.h5')

lstm_5.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])
```

Evaluate Models

```

In [59]: # Baseline
print(baseline.evaluate(X_baseline, y))

# LSTM
print(lstm.evaluate(X_lstm, y))

# One vs. All
one_star_ps = lstm_1.predict(X_lstm)
two_star_ps = lstm_2.predict(X_lstm)
three_star_ps = lstm_3.predict(X_lstm)
four_star_ps = lstm_4.predict(X_lstm)
five_star_ps = lstm_5.predict(X_lstm)

data = [one_star_ps.flatten(), two_star_ps.flatten(), three_star_ps.flatten(),
four_star_ps.flatten(), five_star_ps.flatten()]
cols = [1, 2, 3, 4, 5]
ps = pd.DataFrame(data=data, index=cols).T

ps["ova_pred"] = ps.idxmax(axis=1)

print([MAE(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1)),
Accuracy(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1))])

500/500 [=====] - 0s 86us/step
[2.0396031694412233, 0.2720000147819519]
500/500 [=====] - 1s 1ms/step
[1.3660213603973388, 0.40400001406669617]
[0.988, 0.094]

```

Attempt Ensemble

```

In [60]: # Baseline
baseline_preds = pd.DataFrame(baseline.predict(X_baseline), columns=cols)
baseline_preds['baseline_pred'] = baseline_preds.idxmax(axis=1)

# LSTM
lstm_preds = pd.DataFrame(lstm.predict(X_lstm), columns=cols)
lstm_preds['lstm_pred'] = lstm_preds.idxmax(axis=1)

# One vs. all
ova_preds = ps

all_preds = pd.DataFrame([baseline_preds['baseline_pred'], lstm_preds['lstm_pr
ed'], ova_preds['ova_pred']]).T
all_preds["final_pred"] = all_preds.mode(axis=1)[0]

print([MAE(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols).idxmax(
axis=1)), Accuracy(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols)
.idxmax(axis=1))])


[0.844, 0.246]

```

Misc.

```
In [61]: sns.countplot(all_preds["final_pred"])
```

```
Out[61]: <matplotlib.axes._subplots.AxesSubplot at 0x239868a9148>
```

**Challenge 6**

```
In [62]: c6 = pd.read_json("./yelp_challenge_6_with_answers.jsonl", lines = True)
print(c6.shape)
c6.head()
```

```
(500, 3)
```


```
Out[62]:
```

	review_id	text	stars
0	60	Amazing for Trees\n\n\$20 for a 5 gallon . I wi...	5
1	61	How the hell can Taco Bell be closed before mi...	5
2	62	I actually had no intention of visiting this p...	5
3	63	Yesterday around 3:30 pm I was driving west on...	5
4	64	DR FITZMAURICE did surgery on both hands on th...	5

Quick EDA

```
In [63]: sns.countplot(c6['stars'])
```

```
Out[63]: <matplotlib.axes._subplots.AxesSubplot at 0x237b1342a48>
```



Pre-processing

```
In [64]: c6['text'] = c6['text'].apply(clean_text)
c6.head()
```

```
Out[64]:
```

	review_id	text	stars
0	60	amaz for tree 20 for a 5 gallon i will never g...	5
1	61	how the hell can taco bell be close befor midn...	5
2	62	i actual had no intent of visit thi place at a...	5
3	63	yesterday around 3 30 pm i wa drive west on pi...	5
4	64	dr fitzmauric did surgeri on both hand on the ...	5

Load previous tokenizer

```
In [65]: X = c6['text'].fillna('').values
y = pd.get_dummies(c6['stars'])

# with open('tokenizer.pickle', 'rb') as handle:
#     tokenizer = pickle.load(handle)

max_words

necc_cols = [1, 2, 3, 4, 5]
for col in necc_cols:
    if col not in y.columns:
        y[col] = 0

y = y[necc_cols]
y = y.values

X_baseline = tokenizer.texts_to_matrix(X)
X_lstm = tokenizer.texts_to_sequences(X)
X_lstm = pad_sequences(X_lstm, maxlen=400)
```

Load and compile models

```
In [ ]: # Baseline
baseline = load_model('./models/baseline.h5')

baseline.compile(loss='categorical_crossentropy',
                  optimizer=optimizer,
                  metrics=['accuracy'])

# LSTM
lstm = load_model('./models/lstm.h5')

lstm.compile(loss='categorical_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])

# One vs. all
lstm_1 = load_model('./models/one_star.h5')

lstm_1.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_2 = load_model('./models/two_star.h5')

lstm_2.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_3 = load_model('./models/three_star.h5')

lstm_3.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_4 = load_model('./models/four_star.h5')

lstm_4.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_5 = load_model('./models/five_star.h5')

lstm_5.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])
```

Evaluate Models


```

In [66]: # Baseline
print(baseline.evaluate(X_baseline, y))

# LSTM
print(lstm.evaluate(X_lstm, y))

# One vs. All
one_star_ps = lstm_1.predict(X_lstm)
two_star_ps = lstm_2.predict(X_lstm)
three_star_ps = lstm_3.predict(X_lstm)
four_star_ps = lstm_4.predict(X_lstm)
five_star_ps = lstm_5.predict(X_lstm)

data = [one_star_ps.flatten(), two_star_ps.flatten(), three_star_ps.flatten(),
four_star_ps.flatten(), five_star_ps.flatten()]
cols = [1, 2, 3, 4, 5]
ps = pd.DataFrame(data=data, index=cols).T

ps["ova_pred"] = ps.idxmax(axis=1)

print([MAE(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1)),
Accuracy(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1))])

500/500 [=====] - 0s 100us/step
[2.3629399967193603, 0.4339999854564667]
500/500 [=====] - 0s 530us/step
[2.310380277633667, 0.44200000166893005]
[2.196, 0.446]

```

Attempt Ensemble

```

In [67]: # Baseline
baseline_preds = pd.DataFrame(baseline.predict(X_baseline), columns=cols)
baseline_preds['baseline_pred'] = baseline_preds.idxmax(axis=1)

# LSTM
lstm_preds = pd.DataFrame(lstm.predict(X_lstm), columns=cols)
lstm_preds['lstm_pred'] = lstm_preds.idxmax(axis=1)

# One vs. all
ova_preds = ps

all_preds = pd.DataFrame([baseline_preds['baseline_pred'], lstm_preds['lstm_pr
ed'], ova_preds['ova_pred']]).T
all_preds["final_pred"] = all_preds.mode(axis=1)[0]

print([MAE(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols).idxmax(
axis=1)), Accuracy(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols)
.idxmax(axis=1))])


[2.07, 0.464]

```

Misc.

```
In [68]: sns.countplot(all_preds["final_pred"])
```

```
Out[68]: <matplotlib.axes._subplots.AxesSubplot at 0x237b13b30c8>
```

**Challenge 3**

```
In [69]: c3 = pd.read_json("./yelp_challenge_3_with_answers.jsonl", lines = True)
print(c3.shape)
c3.head()
```

```
(534, 3)
```


```
Out[69]:
```

	review_id	text	stars
0	30	We stopped here for lunch today and were pleas...	4
1	31	We went for a quick lunch here - it's all reas...	3
2	32	Very bad food, avoid it. We were a group of 4 ...	2
3	33	Bring a friend or two to help open the door. I...	3
4	34	Ukai serves some of the best sushi and sashimi...	4

Quick EDA

```
In [70]: sns.countplot(c3['stars'])
```

```
Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x237b138d948>
```



Pre-processing

```
In [71]: c3['text'] = c3['text'].apply(clean_text)
c3.head()
```

```
Out[71]:
```

	review_id	text	stars
0	30	we stop here for lunch today and were pleasant...	4
1	31	we went for a quick lunch here it s all reason...	3
2	32	veri bad food avoid it we were a group of 4 an...	2
3	33	bring a friend or two to help open the door i ...	3
4	34	ukai serv some of the best sushi and sashimi i...	4

Load previous tokenizer

```
In [72]: X = c3['text'].fillna('').values
y = pd.get_dummies(c3['stars'])

# with open('tokenizer.pickle', 'rb') as handle:
#     tokenizer = pickle.load(handle)

max_words

necc_cols = [1, 2, 3, 4, 5]
for col in necc_cols:
    if col not in y.columns:
        y[col] = 0

y = y[necc_cols]
y = y.values

X_baseline = tokenizer.texts_to_matrix(X)
X_lstm = tokenizer.texts_to_sequences(X)
X_lstm = pad_sequences(X_lstm, maxlen=400)
```

Load and compile models

```
In [ ]: # Baseline
baseline = load_model('./models/baseline.h5')

baseline.compile(loss='categorical_crossentropy',
                 optimizer=optimizer,
                 metrics=['accuracy'])

# LSTM
lstm = load_model('./models/lstm.h5')

lstm.compile(loss='categorical_crossentropy',
             optimizer=optimizer,
             metrics=['accuracy'])

# One vs. all
lstm_1 = load_model('./models/one_star.h5')

lstm_1.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_2 = load_model('./models/two_star.h5')

lstm_2.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_3 = load_model('./models/three_star.h5')

lstm_3.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_4 = load_model('./models/four_star.h5')

lstm_4.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])

lstm_5 = load_model('./models/five_star.h5')

lstm_5.compile(loss='binary_crossentropy',
               optimizer=optimizer,
               metrics=['accuracy'])
```

Evaluate Models

```

In [73]: # Baseline
print(baseline.evaluate(X_baseline, y))

# LSTM
print(lstm.evaluate(X_lstm, y))

# One vs. All
one_star_ps = lstm_1.predict(X_lstm)
two_star_ps = lstm_2.predict(X_lstm)
three_star_ps = lstm_3.predict(X_lstm)
four_star_ps = lstm_4.predict(X_lstm)
five_star_ps = lstm_5.predict(X_lstm)

data = [one_star_ps.flatten(), two_star_ps.flatten(), three_star_ps.flatten(),
four_star_ps.flatten(), five_star_ps.flatten()]
cols = [1, 2, 3, 4, 5]
ps = pd.DataFrame(data=data, index=cols).T

ps["ova_pred"] = ps.idxmax(axis=1)

print([MAE(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1)),
Accuracy(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1))])

```

```

534/534 [=====] - 0s 92us/step
[1.1802164695682598, 0.567415714263916]
534/534 [=====] - 0s 562us/step
[0.91535808955239, 0.6086142063140869]
[0.5823970037453183, 0.548689138576779]

```

Attempt Ensemble

```

In [74]: # Baseline
baseline_preds = pd.DataFrame(baseline.predict(X_baseline), columns=cols)
baseline_preds['baseline_pred'] = baseline_preds.idxmax(axis=1)

# LSTM
lstm_preds = pd.DataFrame(lstm.predict(X_lstm), columns=cols)
lstm_preds['lstm_pred'] = lstm_preds.idxmax(axis=1)

# One vs. all
ova_preds = ps

all_preds = pd.DataFrame([baseline_preds['baseline_pred'], lstm_preds['lstm_pr
ed'], ova_preds['ova_pred']]).T
all_preds["final_pred"] = all_preds.mode(axis=1)[0]

print([MAE(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols).idxmax(
axis=1)), Accuracy(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols)
.idxmax(axis=1))])

```

```


[0.5074906367041199, 0.599250936329588]

```

Misc.

```
In [75]: sns.countplot(all_preds["final_pred"])
```

```
Out[75]: <matplotlib.axes._subplots.AxesSubplot at 0x237b14dde48>
```

**Challenge 8**

```
In [76]: c8 = pd.read_json("./yelp_challenge_8_with_answers.jsonl", lines = True)
print(c8.shape)
c8.head()
```

```
(500, 3)
```


```
Out[76]:
```

	review_id	text	stars
0	qOOv-A-vo3kMT0yi4jlllg	Not bad for fast food.	4
1	uqxo6B6w_sIDSAGr0k_0A	Une institution du café	4
2	0o_gGSU0m_4QyNLWEHKgug	J ai vraiment aimé !!!!	4
3	BKAj-fKWW5G3yt3xAkbUCQ	They have good poutine.	4
4	fAhp8lwuGNT0ywKmsCs6VQ	Very old and dirty vans.	1

Quick EDA

```
In [77]: sns.countplot(c8['stars'])
```

```
Out[77]: <matplotlib.axes._subplots.AxesSubplot at 0x237b13e2908>
```



Pre-processing

```
In [78]: c8['text'] = c8['text'].apply(clean_text)
c8.head()
```

C:\Users\Tanner\Anaconda3\envs\yelp\lib\site-packages\bs4__init__.py:398: UserWarning: "https://casetext.com/case/united-states-v-butterbaugh-2" looks like a URL. BeautifulSoup is not an HTTP client. You should probably use an HTTP client like requests to get the document behind the URL, and feed that document to BeautifulSoup.

markup

```
Out[78]:
```

	review_id	text	stars
0	qOOv-A-vo3kMT0yi4jlllg	not bad for fast food	4
1	uqxkO6B6w_sIDSAGr0k_0A	une institut du caf	4
2	0o_gGSU0m_4QyNLWEHKgug	j ai vraiment aim	4
3	BKAj-fKWW5G3yt3xAkbUCQ	they have good poutin	4
4	fAhp8lwuGNT0ywKmsCs6VQ	veri old and dirti van	1

Load previous tokenizer


```
In [79]: X = c8['text'].fillna('').values
y = pd.get_dummies(c8['stars'])

# with open('tokenizer.pickle', 'rb') as handle:
#     tokenizer = pickle.load(handle)

max_words

necc_cols = [1, 2, 3, 4, 5]
for col in necc_cols:
    if col not in y.columns:
        y[col] = 0

y = y[necc_cols]
y = y.values

X_baseline = tokenizer.texts_to_matrix(X)
X_lstm = tokenizer.texts_to_sequences(X)
X_lstm = pad_sequences(X_lstm, maxlen=400)
```

Load and compile models

```
In [ ]: # Baseline
baseline = load_model('./models/baseline.h5')

baseline.compile(loss='categorical_crossentropy',
                  optimizer=optimizer,
                  metrics=['accuracy'])

# LSTM
lstm = load_model('./models/lstm.h5')

lstm.compile(loss='categorical_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])

# One vs. all
lstm_1 = load_model('./models/one_star.h5')

lstm_1.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_2 = load_model('./models/two_star.h5')

lstm_2.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_3 = load_model('./models/three_star.h5')

lstm_3.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_4 = load_model('./models/four_star.h5')

lstm_4.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])

lstm_5 = load_model('./models/five_star.h5')

lstm_5.compile(loss='binary_crossentropy',
                optimizer=optimizer,
                metrics=['accuracy'])
```

Evaluate Models

```
In [80]: # Baseline
print(baseline.evaluate(X_baseline, y))

# LSTM
print(lstm.evaluate(X_lstm, y))

# One vs. All
one_star_ps = lstm_1.predict(X_lstm)
two_star_ps = lstm_2.predict(X_lstm)
three_star_ps = lstm_3.predict(X_lstm)
four_star_ps = lstm_4.predict(X_lstm)
five_star_ps = lstm_5.predict(X_lstm)

data = [one_star_ps.flatten(), two_star_ps.flatten(), three_star_ps.flatten(),
four_star_ps.flatten(), five_star_ps.flatten()]
cols = [1, 2, 3, 4, 5]
ps = pd.DataFrame(data=data, index=cols).T

ps["ova_pred"] = ps.idxmax(axis=1)

print([MAE(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1)),
Accuracy(ps["ova_pred"], pd.DataFrame(data=y, columns=cols).idxmax(axis=1))])

500/500 [=====] - 0s 84us/step
[1.0328141555786132, 0.6439999938011169]
500/500 [=====] - 0s 534us/step
[0.928207818031311, 0.6359999775886536]
[0.848, 0.574]
```

Attempt Ensemble

```
In [81]: # Baseline
baseline_preds = pd.DataFrame(baseline.predict(X_baseline), columns=cols)
baseline_preds['baseline_pred'] = baseline_preds.idxmax(axis=1)

# LSTM
lstm_preds = pd.DataFrame(lstm.predict(X_lstm), columns=cols)
lstm_preds['lstm_pred'] = lstm_preds.idxmax(axis=1)

# One vs. all
ova_preds = ps

all_preds = pd.DataFrame([baseline_preds['baseline_pred'], lstm_preds['lstm_pr
ed'], ova_preds['ova_pred']]).T
all_preds["final_pred"] = all_preds.mode(axis=1)[0]

print([MAE(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols).idxmax(
axis=1)), Accuracy(all_preds["final_pred"], pd.DataFrame(data=y, columns=cols)
.idxmax(axis=1))])

[0.624, 0.626]
```

Misc.

```
In [82]: sns.countplot(all_preds["final_pred"])
```

```
Out[82]: <matplotlib.axes._subplots.AxesSubplot at 0x237c49a9848>
```

